


GPU compute and memory scaling

Each series is normalized to its A100 value · A100 80GB SXM = 1.0



Compute: dense BF16 / FP16 Tensor Core peak. Capacity and bandwidth are separate series.
HBM and L2: per GPU. L1: combined L1 / shared-memory / texture pool per SM, not dedicated L1 cache.
* B100: preliminary specs; L1 uses the Blackwell architecture value. Missing L2 data are not interpolated.
Sources: NVIDIA product specifications, architecture/tuning guides and HGX reference architecture; see accompanying data and notes.